

Cifar-10 dataset - image classification

The **CIFAR-10 dataset** consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

0-airplane

1-automobile

2-bird

3-cat

4-deer

5-dog

6-frog

7-horse

8-ship

9-truck

1.Import libraries

```
In [1]: import pickle
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from keras.models import Sequential
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense, Dropout
from keras import regularizers
from keras.optimizers import SGD
from keras.preprocessing.image import ImageDataGenerator
from keras.utils import np_utils
import keras

Using TensorFlow backend.
```

2.Load data

```
In [2]: def load_train_data(n):
with open('test_batch_'+str(n)+'.npy') as file:
    batch = pickle.load(file, encoding='latin1')
    features = batch['data']
    Target = batch['labels']
    return features, Target

In [3]: batch_1, Target_1 = load_train_data(1)
batch_2, Target_2 = load_train_data(2)
batch_3, Target_3 = load_train_data(3)
batch_4, Target_4 = load_train_data(4)
batch_5, Target_5 = load_train_data(5)

In [4]: print('train batch data shape are ',batch_1.shape,batch_2.shape,batch_3.shape,batch_4.shape,batch_5.shape)
print('train label data shape are ',len(Target_1),len(Target_2),len(Target_3),len(Target_4),len(Target_5))

train batch data shape are (10000, 3072) (10000, 3072) (10000, 3072) (10000, 3072) (10000, 3072)
train label data shape are 10000 10000 10000 10000 10000
```

```
In [5]: with open('test_batch_', 'rb') as file:
    batch = pickle.load(file, encoding='latin1')
    X_test = batch['data']
    Y_test = batch['labels']
    print('test batch data and label data shape are ',X_test.shape,len(Y_test))

test batch data and label data shape are (10000, 3072) 10000
```

3.Image Plotting

```
In [6]: def plot_image(batch,labels,num):
    images =batch.reshape((len(batch), 3, 32, 32)).transpose(0,2,3,1)
    dict = {'0':'Airplane', 1:'Automobile', 2:'Bird', 3:'Cat', 4:'Deer', 5:'Dog', 6:'Frog', 7:'Horse', 8:'Ship', 9:'Truck'}
    fig = plt.figure(figsize = (2,2))
    plt.imshow(images[num])
    plt.title(dict[labels[num]])

In [7]: plot_image(batch_1,Target_1,100)

In [8]: plot_image(batch_2,Target_2,100)

In [9]: plot_image(batch_3,Target_3,100)

In [10]: plot_image(batch_4,Target_4,100)

In [11]: plot_image(batch_5,Target_5,100)
```

4: Data Prepatation

```
In [12]: X_train = np.append(batch_1, batch_2,axis=0)
X_train = np.append(X_train, batch_3,axis=0)
X_train = np.append(X_train, batch_4,axis=0)
X_train = np.append(X_train, batch_5,axis=0)
Y_train = np.append(Target_1, Target_2,axis=0)
Y_train = np.append(Y_train, Target_3,axis=0)
Y_train = np.append(Y_train, Target_4,axis=0)
Y_train = np.append(Y_train, Target_5,axis=0)

In [13]: print(X_train.shape,len(Y_train))
(50000, 3072) 50000

In [14]: X_train = X_train.reshape((len(X_train), 3, 32, 32)).transpose(0,2,3,1)
Y_train = np_utils.to_categorical(Y_train, 10)

In [15]: X_test = X_test.reshape((len(X_test), 3, 32, 32)).transpose(0,2,3,1)
Y_test = np_utils.to_categorical(Y_test, 10)

In [16]: print(X_train.shape,X_test.shape,len(Y_train),len(Y_test))
(50000, 32, 32, 3) (10000, 32, 32, 3) 50000 10000

In [17]: X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0

In [ ]:
```

CNN model

Conv2D layer

- keras.layers.Conv2D(filters, kernel_size, strides=(1,1), padding='valid', data_format=None, dilation_rate=(1,1), activation=None, use_bias=True, kernel_initializer='glorot_uniform', bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None, kernel_constraint=None, bias_constraint=None)**

Parameters

1. Filters

The filter value depends on the complexity of the dataset and the depth of neural network.

In this architecture start with 32 filters.

2. Kerner_size.

The second required parameter in the Keras Conv2D class is the kernel_size , a 2-tuple specifying the width and height of the 2D convolution window.

As the size of the image is 32 x 32 will choose either 1x1 or 3x3 kernel size.

3. Strides.

The strides parameter is a 2-tuple of integers, specifying the 'step' of the convolution along the x and y axis of the input volume.

The strides value defaults to (1, 1) as I use the same.

4. Padding.

To preserve the spatial dimensions of the volume, i use "same" for the padding.

5. Activation

The activation parameter helps keep code cleaner and won't have an impact on the performance of Convolutional Neural Network.

1. kernel_initializer

The kernel_initializer defaults to glorot_uniform , the Xavier Glorot uniform initialization method, which is perfectly fine for the majority of tasks; however, for deeper neural networks want to use "he_normal" if network has a large number of parameters.

1. The kernel_regularizer , bias_regularizer , and activity_regularizer

Controls the type and amount of regularization method applied to the Conv2D layer.

Applying regularization helps to:

- Reduce the effects of overfitting
- Increase the ability of your model to generalize

1. Rest all parameters are kept as default.

MaxPooling operation for 2D

Max pooling operation for spatial data.

- keras.layers.MaxPooling2D(pool_size=(2, 2), strides=None, padding='valid', data_format=None)**

Model 1

Lets check the performance with below parameters

- Initial filter with 64
- kernel size 3x3
- activation - ReLU
- Kernel_initializer - he_normal
- kernel_regularizer - 0.001
- 5 conv2D layers, 3 maxpooling2D, 2 dense or fully connected layers, one flatten
- optimizer SGD
- loss - cross entropy
- Metrics - accuracy
- batch_size = 32
- epochs - 10
- No cross validation

```
In [18]: # Lets build a model with conv2d layers
model = Sequential()
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=reg
ularizers.l2(0.001), input_shape=(32, 32, 3)))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=reg
ularizers.l2(0.001)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=reg
ularizers.l2(0.001)))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=reg
ularizers.l2(0.001)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=reg
ularizers.l2(0.001)))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(10, activation='softmax'))
model.summary()

WARNING:tensorflow:From C:\Users\Udhayan\Anaconda3\lib\site-packages\tensorflow\python\framework
\ops\op_def_library.py:263: colocate_with (from tensorflow.python.framework.framework_ops) is deprecated a
nd will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 30, 30, 64)	1792
conv2d_2 (Conv2D)	(None, 28, 28, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_3 (Conv2D)	(None, 12, 12, 64)	36928
conv2d_4 (Conv2D)	(None, 10, 10, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_5 (Conv2D)	(None, 3, 3, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten_1 (Flatten)	(None, 64)	0
dense_1 (Dense)	(None, 128)	8320
dense_2 (Dense)	(None, 10)	1290
Total params: 159,114		
Trainable params: 159,114		
Non-trainable params: 0		

```
In [19]: epochs = 10
sgd = SGD(lr=1e-2, momentum=0.9, decay=1e-2/epochs)
model.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(X_train,y_train,epochs=epochs,batch_size = 32)

WARNING:tensorflow:From C:\Users\Udhayan\Anaconda3\lib\site-packages\tensorflow\python\ops\mat
h_ops.py:3966: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed
in a future version.
Instructions for updating:
Use tf.cast instead.

Epoch 1/10
50000/50000 [=====] - 360s 7ms/step - loss: 2.2050 - acc: 0.3811
Epoch 2/10
50000/50000 [=====] - 117s 2ms/step - loss: 1.6950 - acc: 0.5551
Epoch 3/10
50000/50000 [=====] - 259s 5ms/step - loss: 1.4814 - acc: 0.6193 5s - 1
cost: 1.4828 - acc = ETA: 3s - loss: 1.4827 - ETA: 1s - loss: 1.4822 -
Epoch 4/10
50000/50000 [=====] - 260s 5ms/step - loss: 1.3625 - acc: 0.6529 5s - 1
cost: 1.3638 - acc: 0 - ETA:
Epoch 5/10
50000/50000 [=====] - 264s 5ms/step - loss: 1.2680 - acc: 0.6814
Epoch 6/10
50000/50000 [=====] - 269s 5ms/step - loss: 1.1962 - acc: 0.7037
Epoch 7/10
50000/50000 [=====] - 277s 6ms/step - loss: 1.1463 - acc: 0.7171
Epoch 8/10
50000/50000 [=====] - 307s 6ms/step - loss: 1.0942 - acc: 0.7328
Epoch 9/10
50000/50000 [=====] - 250s 5ms/step - loss: 1.0562 - acc: 0.7483
Epoch 10/10
50000/50000 [=====] - 240s 5ms/step - loss: 1.0214 - acc: 0.7552

Out[19]: <keras.callbacks.History at 0x23076581d0>
```

```
In [20]: test_loss,test_acc = model.evaluate(X_test,y_test)
test_acc

10000/10000 [=====] - 18s 2ms/step

Out[20]: 0.7106
```

Observation

The train accuracy for the above model with 10 epochs is 75.5% The test accuracy is 71% so if the epochs increased accuracy will vary

The difference in accuracy for train and test is approx 4% it may change with other unknown pictures

model 2

Lets check the performance by changing initial filter in increasing order with powers of 2

- Initial filter with 32
- kernel size 3x3
- activation - ReLU
- Kernel_initializer - he_normal
- kernel_regularizer - 0.001
- 5 conv2D layers, 3 maxpooling2D, 2 dense or fully connected layers, one flatten
- optimizer SGD
- loss - cross entropy
- Metrics - accuracy
- batch_size = 32
- epochs - 10
- No cross validation

```
In [22]: model2 = Sequential()
model2.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001), input_shape=(32, 32, 3)))
model2.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model2.add(MaxPooling2D((2, 2)))
model2.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model2.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model2.add(MaxPooling2D((2, 2)))
model2.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model2.add(MaxPooling2D((2, 2)))
model2.add(Flatten())
model2.add(Dense(128, activation='relu'))
model2.add(Dense(10, activation='softmax'))
model2.summary()
```

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 30, 30, 32)	896
conv2d_7 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d_4 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_8 (Conv2D)	(None, 12, 12, 64)	18496
conv2d_9 (Conv2D)	(None, 10, 10, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_10 (Conv2D)	(None, 3, 3, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 1, 1, 128)	0
flatten_2 (Flatten)	(None, 128)	0
dense_3 (Dense)	(None, 128)	16512
dense_4 (Dense)	(None, 10)	1290
Total params: 157,226		
Trainable params: 157,226		
Non-trainable params: 0		

```
In [23]: epochs = 10
sgd = SGD(lr=1e-2, momentum=0.9, decay=1e-2/epochs)
model2.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
model2.fit(X_train,y_train,epochs=epochs,batch_size = 32)

Epoch 1/10
50000/50000 [=====] - 103s 2ms/step - loss: 2.2591 - acc: 0.3634
Epoch 2/10
50000/50000 [=====] - 117s 2ms/step - loss: 1.7776 - acc: 0.5163
Epoch 3/10
50000/50000 [=====] - 116s 2ms/step - loss: 1.5704 - acc: 0.5816
Epoch 4/10
50000/50000 [=====] - 126s 3ms/step - loss: 1.4424 - acc: 0.6203
Epoch 5/10
50000/50000 [=====] - 131s 3ms/step - loss: 1.3423 - acc: 0.6529
Epoch 6/10
50000/50000 [=====] - 136s 3ms/step - loss: 1.2669 - acc: 0.6765
Epoch 7/10
50000/50000 [=====] - 129s 3ms/step - loss: 1.2028 - acc: 0.6969 2s - 1
cost: 1.2023 - - ETA: 1s - loss
Epoch 8/10
50000/50000 [=====] - 129s 3ms/step - loss: 1.1511 - acc: 0.7121 0s - 1
cost: 1.1509 - acc
Epoch 9/10
50000/50000 [=====] - 130s 3ms/step - loss: 1.1018 - acc: 0.7280
Epoch 10/10
50000/50000 [=====] - 133s 3ms/step - loss: 1.0646 - acc: 0.7385

Out[23]: <keras.callbacks.History at 0x2381c90cc0>
```

```
In [24]: test_loss,test_acc = model2.evaluate(X_test,y_test)
test_acc

10000/10000 [=====] - 11s 1ms/step

Out[24]: 0.6629
```

Observation

The above model with filters changing in increasing order decreased the accuracy by 2% and with less computation model 2 reduced test accuracy 1% data from model 1

Model 3

Change in kernel shape 1 X 1

```
In [ ]:
```

```
In [25]: model3 = Sequential()
model3.add(Conv2D(32, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001), input_shape=(32, 32, 3)))
model3.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model3.add(MaxPooling2D((2, 2)))
model3.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model3.add(MaxPooling2D((2, 2)))
model3.add(Conv2D(128, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model3.add(MaxPooling2D((2, 2)))
model3.add(Flatten())
model3.add(Dense(128, activation='relu'))
model3.add(Dense(10, activation='softmax'))
model3.summary()
```

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 32, 32, 32)	128
conv2d_12 (Conv2D)	(None, 32, 32, 32)	1056
max_pooling2d_7 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_13 (Conv2D)	(None, 16, 16, 64)	2112
conv2d_14 (Conv2D)	(None, 16, 16, 64)	4160
max_pooling2d_8 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_15 (Conv2D)	(None, 8, 8, 128)	8320
max_pooling2d_9 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_3 (Flatten)	(None, 2048)	0
dense_5 (Dense)	(None, 128)	262272
dense_6 (Dense)	(None, 10)	1290
Total params: 279,338		
Trainable params: 149,386		
Non-trainable params: 0		

```
In [26]: epochs = 10
sgd = SGD(lr=1e-2, momentum=0.9, decay=1e-2/epochs)
model4.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
model4.fit(X_train,y_train,epochs=epochs,batch_size = 32)

Epoch 1/10
50000/50000 [=====] - 71s 1ms/step - loss: 2.3842 - acc: 0.3185
Epoch 2/10
50000/50000 [=====] - 69s 1ms/step - loss: 1.9200 - acc: 0.4613
Epoch 3/10
50000/50000 [=====] - 68s 1ms/step - loss: 1.7791 - acc: 0.4967
Epoch 4/10
50000/50000 [=====] - 68s 1ms/step - loss: 1.6963 - acc: 0.5162
Epoch 5/10
50000/50000 [=====] - 57s 1ms/step - loss: 1.6320 - acc: 0.5366
Epoch 6/10
50000/50000 [=====] - 54s 1ms/step - loss: 1.5832 - acc: 0.5456
Epoch 7/10
50000/50000 [=====] - 53s 1ms/step - loss: 1.5437 - acc: 0.5575
Epoch 8/10
50000/50000 [=====] - 53s 1ms/step - loss: 1.5086 - acc: 0.5667
Epoch 9/10
50000/50000 [=====] - 53s 1ms/step - loss: 1.4801 - acc: 0.5759
Epoch 10/10
50000/50000 [=====] - 53s 1ms/step - loss: 1.4554 - acc: 0.5816

Out[26]: <keras.callbacks.History at 0x2381cca0a20>
```

```
In [27]: test_loss,test_acc = model3.evaluate(X_test,y_test)
test_acc

10000/10000 [=====] - 5s 47us/step

Out[27]: 0.5748
```

Observation

By reducing kernel size the train accuracy dropped by almost 17% But the difference between the test and train is only 2% it requires only less computation compare to above models so kernel size 1x1 is not a good choice

Model -4

```
In [28]: model4 = Sequential()
model4.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001), input_shape=(32, 32, 3)))
model4.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model4.add(MaxPooling2D((2, 2)))
model4.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model4.add(MaxPooling2D((2, 2)))
model4.add(Conv2D(64, (1, 1), activation='relu', kernel_initializer='he_normal', kernel_regularizer=re
gularizers.l2(0.001)))
model4.add(MaxPooling2D((2, 2)))
model4.add(Flatten())
model4.add(Dense(128, activation='relu'))
model4.add(Dense(10, activation='softmax'))
model4.summary()
```

Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 32, 32, 64)	256
conv2d_17 (Conv2D)	(None, 32, 32, 64)	4160
max_pooling2d_10 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_18 (Conv2D)	(None, 16, 16, 64)	4160
conv2d_19 (Conv2D)	(None, 16, 16, 64)	4160
max_pooling2d_11 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_20 (Conv2D)	(None, 8, 8, 128)	4160
max_pooling2d_12 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_4 (Flatten)	(None, 1024)	0
dense_7 (Dense)	(None, 128)	131200
dense_8 (Dense)	(None, 10)	1290
Total params: 149,386		
Trainable params: 149,386		
Non-trainable params: 0		

```
In [29]: epochs = 10
sgd = SGD(lr=1e-2, momentum=0.9, decay=1e-2/epochs)
model4.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
model4.fit(X_train,y_train,epochs=epochs,batch_size = 32)

Epoch 1/10
50000/50000 [=====] - 78s 2ms/step - loss: 2.2423 - acc: 0.3753
Epoch 2/10
50000/50000 [=====] - 78s 2ms/step - loss: 1.8545 - acc: 0.4848
Epoch 3/10
50000/50000 [=====] - 76s 2ms/step - loss: 1.7188 - acc: 0.5180
Epoch 4/10
50000/50000 [=====] - 77s 2ms/step - loss: 1.6352 - acc: 0.5387
Epoch 5/10
50000/50000 [=====] - 77s 2ms/step - loss: 1.5727 - acc: 0.5562
Epoch 6/10
50000/50000 [=====] - 77s 2ms/step - loss: 1.5296 - acc: 0.5760
Epoch 7/10
50000/50000 [=====] - 76s 2ms/step - loss: 1.4938 - acc: 0.5761
Epoch 8/10
50000/50000 [=====] - 76s 2ms/step - loss: 1.4638 - acc: 0.5816
Epoch 9/10
50000/50000 [=====] - 78s 2ms/step - loss: 1.4187 - acc: 0.5866
Epoch 10/10
50000/50000 [=====] - 78s 2ms/step - loss: 1.4187 - acc: 0.5866
```